
On the Value of Abstractions: Abstraction Selection in Bounded Program Synthesis

Adam Cuculich

adam@cuculich.me

Abstract

Reusable abstractions can shorten programs in program synthesis. However, each item also expands a solver’s search space. With a fixed budget, this added search can prevent the solver from reaching a shorter solution. We ask when adding reusable abstractions improves bounded synthesis and what determines their value. The answer matters because future savings must justify selection and added search costs. To study this question, we compare two abstraction-selection procedures in Pattern Builder, a synthetic grid benchmark. The benchmark uses deterministic, size-ordered search. Across 30 random problem worlds, both procedures choose from the same set of 50 abstraction candidates. These candidates are composite programs found while solving starter problems. A selected candidate becomes a reusable item available to the solver. Compression-based selection favors candidates that shorten solutions found for selected problems. Validation Search Utility favors candidates that reduce capped search cost on separate validation problems. Together, the solver’s primitive items and selected abstractions form its library. We evaluate each resulting library on the same held-out test problems. We then vary the number of selected abstractions and the search budget.

With ten abstractions, both procedures outperformed the original domain-specific language (DSL), which contained no added abstractions. Validation Search Utility solved 4.08 more test problems per 100 than compression-based selection (95 percent CI 2.46 to 5.71). For this comparison, compression-based selection used solutions acquired from the same validation problems. The difference includes the selection criterion and which validation problems produced solutions. Validation Search Utility still showed an advantage over compression-based selection using solutions from 100 starter problems, but the observed advantage was statistically uncertain. At 30,000 attempts, mean performance peaked near 11 abstractions and fell below the original DSL at 20. For both procedures, performance also decreased when the library grew from one selected abstraction to two. To test whether the budget caused this decrease, we held those libraries fixed and increased only the budget. The decrease reversed. This result supports a search-cutoff explanation. Added abstractions can consume the budget before the solver reaches larger programs where they help. Validation Search Utility also required more search during selection, and some cases did not recover that cost. Therefore, solution compression alone does not measure an abstraction’s value under bounded search. A library should be evaluated with its intended solver, budget, workload, and expected reuse.

1 Introduction

Reusable abstractions can accelerate program synthesis. They can replace several operations with one library item. However, each abstraction also increases the set of programs that a solver can construct. With a fixed search budget, the number and content of the abstractions can determine whether the solver finds a solution. This paper asks which abstractions to keep, how many to keep,

and when the search budget prevents a useful library from helping. We study these questions in Pattern Builder. The solver must reproduce a binary 10×10 target grid without an input grid. It combines six primitive grids with fixed unary and binary operators until an expression matches the target. Figure 1 shows one solution. It uses `add`, which returns the union of two grids.




Figure 1: A grid target and one expression that reproduces it.

The solver can store an intermediate grid as a library item. We call this stored grid an *abstraction*. In Figure 1, the solver could store the union of the vertical line and the square. This would reduce the solution from two operations to one. The abstraction is a fixed grid, not a learned operator. We change the library and keep the operators fixed. An abstraction can shorten a solution, but it also becomes an available argument for the operators. This creates more small programs. The solver examines all smaller programs first. These programs can use the budget before the solver reaches a larger program where the abstraction helps. Therefore,

shorter known solutions $\nRightarrow$ less search on unseen problems.

Compression selects abstractions that remove the most operations from known solutions. It measures the size of the shortcut each abstraction creates by making those solutions expressible with fewer operations. Its score excludes the extra search that the larger library creates. Validation search utility selects abstractions by their effect on bounded search for separate validation targets. All procedures use the same candidate menu. It contains composite programs that can help solve at least two starter problems. Validation data can change which candidates are selected, but cannot add candidates. Test data only evaluate the selected libraries on unseen targets. Both procedures use the same 25 validation problems. Validation Search Utility scores all search outcomes, including failures, while compression scores only the solution programs found for those problems and omits unsolved problems. **Experiment I** compares the procedures and tests the effect of similarity between the starter and test problem sets. **Experiment II** varies library size. **Experiment III** varies the search budget while the libraries stay fixed. Together, they test selection method, library size, and search budget.

2 Related Work

Library learning adds reusable program components to make later synthesis easier. Many systems select these components by compressing known solution programs. DreamCoder learns a function library together with a neural search policy (Ellis et al., 2021). Stitch finds functions that minimize the cost of a rewritten program corpus (Bowers et al., 2023). A recent study used Pattern Builder to examine human helper choices under changing task structure (Hernandez Cano et al., 2026). Human choices were consistent with compression of expected future solutions. Our study instead measures how automated abstraction selection changes bounded search.

Iba described the underlying search tradeoff. Macros shorten solution paths and increase the search branching factor (Iba, 1989). AbstractBeam later added Stitch abstractions to guided bottom-up synthesis (Zenkner et al., 2024). It improved performance on related tasks. A distribution shift removed this benefit, and unused abstractions could enlarge the language. AMaze provides a close technical comparison. It modifies a DSL using predicted synthesis time. AMaze predicts this time

from the enumeration rank of solver-specific program fragments. Its whole-program compression baseline slowed two of four synthesizers (Ye et al., 2026). Our work measures capped cost directly with exact enumeration over a fixed menu within each seed. It varies library size and changes only the budget for fixed libraries. It also measures the work used to select each library. These controls test how selection method, library size, and budget jointly determine abstraction value.

3 Methodology

3.1 Experiment Overview

We test how abstraction selection affects bounded-search performance. Later experiments change library size and search budget. Experiment I compares four selection procedures. Each procedure selects from the same candidate menu. We then test each selected library on held-out problems. Starter problems create the candidate menu and supply solutions for starter-based compression. The validation-based procedures use separate validation problems. We reserve the test problems for held-out evaluation. We repeat the experiment across 30 random problem worlds. A seed creates one world with 12 hidden motifs and one starter set. Six primary conditions share that world. Each condition has separate validation and test problems. Table 1 gives the size and purpose of each problem set.

Table 1: Problem sets and their use.

Problem Set	Size	Purpose
Starter Problems	100	Creates the candidate menu and supplies problems for starter-based compression.
Validation Problems	25	Guides validation-solution compression and validation utility.
Test Problems	100	Measures each selected library on held-out problems.

Each seed has one menu of 50 candidate abstractions. Every selection procedure within that seed uses this menu. Solutions to starter problems create the menu. Validation evidence changes which candidates a procedure selects. Test problems only measure the selected libraries. Table 2 defines the four selection procedures.

Table 2: Four selection procedures in Experiment I.

ID	Procedure	Problems	Search during selection	Score
S_{100}	Standard compression	100 starter	Acquire solutions; 30,000 attempts per problem	Remaining operations in acquired solutions
S_{25}	Starter-subset compression	25 starter	Acquire solutions; 30,000 attempts per problem	Remaining operations in acquired solutions
M	Validation-solution compression	25 validation	Acquire more solutions; 90,000 attempts per problem	Remaining operations in acquired solutions
U	Validation utility	25 validation	Score each trial library; 30,000 attempts	Mean capped search cost for all 25 targets

Each compression procedure runs a separate target-directed search for each scoring problem. It keeps the first solution found and omits unsolved problems. Utility runs one size-ordered search for each trial library. The resulting output record scores all 25 validation targets. An unreached target receives the maximum cost of 30,000. Hidden generator programs are never selector inputs. A fixed seeded shuffle selects the 25 starter problems for S_{25} . All six conditions within a seed use the same subset. Across the 180 primary Experiment I cells, S_{100} acquired solutions for a mean of 62.70 of 100 problems. S_{25} acquired solutions for 16.33 of 25, and M acquired solutions for 16.59 of 25.

In Experiment I, each procedure selects ten abstractions. We evaluate every selected library on the same 100 test problems with a 30,000-attempt limit. The primary comparison starts with the same 25 validation problem identities. Utility scores all 25 target outcomes. Validation-solution compression scores only acquired solution programs and omits a mean of 8.41 unsolved problems. The comparison holds the candidate menu, test problems, library size, and test budget constant. It includes the scoring objective and the usable evidence. Thus, it compares complete procedures.

3.2 Synthesis and Bounded Search

Each problem supplies one binary target grid T . The primitive-only library L_0 contains six grids: blank, horizontal line, vertical line, diagonal, square, and triangle. The current ordered library is L . A program is a library item or an operator applied to one or two programs. The unary operators are invert, a left-to-right flip, a top-to-bottom flip, and diagonal reflection. The diagonal runs from the upper-left corner to the lower-right corner. Invert changes each black cell to white and each white cell to black. The binary operators are `add` (union), `subtract` (set difference), and `overlap` (intersection). Program P produces grid $\text{eval}(P)$. It solves T when $\text{eval}(P) = T$. Program size is the number of operator applications. A primitive or stored abstraction has size zero. Let B be the maximum number of candidate programs that the solver can try. The deterministic bottom-up solver tries programs in order of increasing size. It stops after B attempts or after it has tried all programs with seven or fewer operator applications. The solver keeps the first program for each distinct output grid. The implementation fixes the order of operators, child-program sizes, and commutative arguments. Each complete program counts as one attempt. This count includes a duplicate that the solver later discards. For utility scoring and final evaluation, one size-ordered search records all outputs for a library. This record gives the first attempt that produced each target. Compression uses separate target-directed searches to acquire solution programs. For target T and library L , the capped search cost is

$$N_B(T, L) = \begin{cases} n, & \text{if target } T \text{ first appears at attempt } n \leq B, \\ B, & \text{if the solver does not reach } T \text{ within the budget.} \end{cases} \quad (1)$$

Unless stated otherwise, $B = 30,000$. The primary outcome is the number of the 100 held-out test targets that the solver reaches.

3.3 The Abstraction Candidate Menu

An abstraction stores a reached grid as a size-zero library item. For each seed, the primitive-only solver tries 30,000 programs and records distinct grids. It runs to the budget. We keep nonblank, nonprimitive grids whose first program has one to four operations. A grid supports a starter target when one invert or reflection produces that target. It also supports the target when one add or subtract with another reached grid produces it. We keep grids that support at least two starter targets. We rank candidates by decreasing support count. We break ties by the attempt at which the solver first reached the grid, then by program string. The top 50 form the candidate menu $\mathcal{C}$, from which abstractions are later selected. Abstraction candidate discovery does not use validation or test data.

3.4 Selection Procedures

Let K be the number of abstractions to select. We run each selection procedure separately. For one run, let $A_0 = ()$. After k selections, $A_k = (a_1, \dots, a_k)$ is the ordered list of selected abstractions. Library $L(A_k)$ contains the primitives followed by this list. The order matters because it changes the solver's search order. The notation $A \parallel a$ means list A followed by abstraction a . For cost function J , the greedy selection procedure is

$$\begin{aligned} a_k &= \arg \min_{\substack{a \in \mathcal{C} \\ a \notin A_{k-1}}} J(A_{k-1} \parallel a), \\ A_k &= A_{k-1} \parallel a_k, \quad k = 1, \dots, K. \end{aligned} \quad (2)$$

Compression and utility use different costs and tie rules. Compression scores acquired solution programs. Let $\mathcal{P}_{\text{sol}}$ contain the first solution found for each solved scoring problem. For program P , let $\text{children}(P)$ be its ordered list of immediate child programs. For selected abstraction list A , the number of remaining operations is

$$n_{\text{ops}}(P, A) = \begin{cases} 0, & \text{if } \text{children}(P) = () \text{ or } \text{eval}(P) \in A, \\ 1 + \sum_{P' \in \text{children}(P)} n_{\text{ops}}(P', A), & \text{otherwise.} \end{cases} \quad (3)$$

A subtree whose output is in A becomes one size-zero library item. Its internal operations contribute zero. Let $\mathcal{T}_{\text{util}}$ contain the targets used to score utility. The two costs are

$$J_{\text{comp}}(A) = \sum_{P \in \mathcal{P}_{\text{sol}}} n_{\text{ops}}(P, A),$$

$$J_{\text{util}}(A) = \frac{1}{|\mathcal{T}_{\text{util}}|} \sum_{T \in \mathcal{T}_{\text{util}}} N_B(T, L(A)). \quad (4)$$

Compression sets $J = J_{\text{comp}}$. Utility sets $J = J_{\text{util}}$. Utility assigns cost B to any target the solver does not reach within B attempts. Thus, when $B = 30,000$, an unreached target has cost 30,000. If two candidates have the same compression cost, Compression compares their text forms in lexicographic order. This rule compares characters from left to right. For example, `add(diagonal, square)` comes before `add(line_vertical, square)` because `d` comes before `l`. If two trial libraries have the same utility cost, Utility first selects the library that solves more targets. If both libraries solve the same number, Utility applies the same text-order rule.

3.5 Problem Sets and Problem-Set Similarity

Each seed defines 12 hidden motifs, with four each at two, three, and four operations. From these motifs, the generator creates 100 starter, 25 validation, and 100 test problems. Each problem uses two or three distinct motifs combined from left to right with uniformly sampled binary operators and up to two unary operators. Fixed schedules set the motif and unary-operator counts. The generator rejects invalid and duplicate grids. Selection procedures receive only target grids or acquired solutions. They do not receive motif identities or hidden programs. Starter problems use motif probabilities $\mathbf{p}_{\text{start}}$. We define $\mathbf{p}_{\text{alt}}$ by reversal or random permutation. Reversal changes $(p_1, \dots, p_{12})$ to $(p_{12}, \dots, p_1)$. Common starter motifs then become rare, and rare starter motifs become common. A random permutation assigns the same probability values to different motifs. The shift parameter $\alpha \in \{0, 0.5, 1\}$ sets the future probabilities:

$$\mathbf{p}(\alpha) = (1 - \alpha)\mathbf{p}_{\text{start}} + \alpha\mathbf{p}_{\text{alt}}.$$

The values $\alpha = 0, 0.5$, and 1 use the starter probabilities, an equal mixture, and the alternative probabilities. In each primary condition, validation and test problems are separate draws from $\mathbf{p}(\alpha)$. Combining two alternatives with three values of α gives six conditions per seed. The two conditions at $\alpha = 0$ use separate target draws. We measure realized problem-set similarity, ρ , between the starter and test sets. It is the Spearman correlation between their motif-count vectors. We use ρ only for analysis. It requires hidden motif identities and the completed test set.

3.6 Experiments

Let q identify a selection procedure, $i \in \{1, \dots, 30\}$ a seed, and $c \in \{1, \dots, 6\}$ a problem-set similarity condition. Outcome $Y_{ic}^q(K, B)$ is the number of test problems that procedure q solves with budget B and its first K abstractions. Thus, $Y_{ic}^q(K, B) \in \{0, \dots, 100\}$. The seed mean and experiment mean are

$$\bar{Y}_i^q(K, B) = \frac{1}{6} \sum_{c=1}^6 Y_{ic}^q(K, B), \quad \hat{\mu}^q(K, B) = \frac{1}{30} \sum_{i=1}^{30} \bar{Y}_i^q(K, B). \quad (5)$$

Each seed gives six related outcomes. We average them into one seed score, then average the 30 scores. Because each seed has six conditions, the experiment mean equals the mean of all 180 cells. Conditions have equal weight. Only the 30 seeds are independent.

Experiment I: Which Abstractions? Experiment I compares the four procedures. It also tests whether problem-set similarity changes the difference between utility and standard compression. Thirty seeds under six primary conditions gave 180 cells. The 30 registered stale-validation cells used validation problems at $\alpha = 0.5$ and test problems at $\alpha = 1$. They formed a separate stress test of outdated validation data. We excluded them from the primary estimates because those estimates require validation and test problems from the same distribution. Each primary cell used $K = 10$, one candidate menu, and the same held-out test targets for all procedures. Table 2 defines U , M , S_{25} , and S_{100} . Standard compression, S_{100} , is the practical comparator. Post hoc, utility also scored the same 25 starter problems as S_{25} . Using the experiment mean $\hat{\mu}^q(K, B)$ defined above, the registered contrasts are

$$\begin{aligned} C_{U-M} &= \hat{\mu}^U(10, 30,000) - \hat{\mu}^M(10, 30,000), \\ C_{M-S_{25}} &= \hat{\mu}^M(10, 30,000) - \hat{\mu}^{S_{25}}(10, 30,000). \end{aligned} \quad (6)$$

A positive value favors the first procedure in each subscript. Contrast C_{U-M} compares the complete validation procedures. Contrast $C_{M-S_{25}}$ compares two complete compression procedures. They differ in scoring-problem source and acquisition budget. Let W_q count, within one seed-condition cell, the candidate-program attempts used to select procedure q 's library. Utility scores $455 = 50 + 49 + \dots + 41$ trial libraries with at most 30,000 attempts each. The work W_M includes validation-solution acquisition. The work $W_{S_{100}}$ treats prior starter-solution acquisition as a sunk cost. These counts exclude compression segmentation. Let $\bar{N}_q$ be the mean capped cost across that cell's 100 held-out test problems. For comparator $q \in \{M, S_{100}\}$, $n_{\text{pay}}(q)$ is the number of future problems needed to recover utility's additional selection work:

$$n_{\text{pay}}(q) = \begin{cases} \left\lceil \frac{\max(0, W_U - W_q)}{\bar{N}_q - \bar{N}_U} \right\rceil, & \text{if } \bar{N}_q > \bar{N}_U, \\ \infty, & \text{if } \bar{N}_q \leq \bar{N}_U. \end{cases} \quad (7)$$

This post hoc calculation assumes constant mean savings and uses only program attempts. We hypothesized that starter-based compression would improve relative to U as starter and test motif frequencies became more similar. Repeated motifs can let starter abstractions help test problems. For assigned shift, we averaged the two alternative assignments. We then calculated each procedure difference at $\alpha = 1$ minus $\alpha = 0$. Regressions related these differences to realized problem-set similarity ρ with seed fixed effects and seed-clustered standard errors.

Experiment II: Library Size. Thirty fresh seeds (6541–6570) and six conditions formed 180 cells. Standard compression, validation-solution compression, and utility each selected one ordered sequence of 20 abstractions. We evaluated every prefix from $K = 0$ to $K = 20$ with $B = 30,000$. The value $K = 0$ is the primitive-only library. Selection ran once for each sequence. For procedures q and q' , set $B = 30,000$. Their performance difference at library size K is

$$D_{q,q'}(K) = \hat{\mu}^q(K, B) - \hat{\mu}^{q'}(K, B). \quad (8)$$

The registered method contrasts were $D_{U,M}(20)$, $D_{U,S_{100}}(20)$, and the change in $D_{U,M}$ from $K = 10$ to $K = 20$. Two registered within-procedure contrasts measured the changes in $\hat{\mu}^U$ and $\hat{\mu}^{S_{100}}$ from $K = 10$ to $K = 20$. Observed peaks and the one-to-two drop were descriptive.

Experiment III: Search Budget. Experiment III reused Experiment II's exact $K = 1$ and $K = 2$ libraries. It varied only $B \in \{30,000, 45,000, 60,000, 90,000\}$. Worlds, targets, abstraction order, size limit, and enumeration order remained fixed. Here, q identifies the fixed library selected in Experiment II. Selection did not run again. For $q \in \{S_{100}, U\}$, the effect of the second abstraction

is $\Delta_q(B)$. The budget interaction is I_q :

$$\begin{aligned}\Delta_q(B) &= \hat{\mu}^q(2, B) - \hat{\mu}^q(1, B), \\ I_q &= \Delta_q(90,000) - \Delta_q(30,000).\end{aligned}\tag{9}$$

A positive I_q means that increasing the search budget raises the $K = 2$ minus $K = 1$ performance difference. We calibrated the budgets on ten previously analyzed sensitivity seeds before we inspected above-30,000 outcomes for the Experiment II seeds.

3.7 Statistical Analysis

Primary comparisons average six related differences between procedures into one seed-level difference. The 30 seed-level differences test whether the advantage repeats across worlds. Similarity analyses keep conditions separate because similarity varies by condition. They treat conditions that share a seed as related. The two primary comparisons use two-sided 95% confidence intervals based on Student’s t distribution. Each interval uses the mean and variation of the 30 differences. The t distribution accounts for uncertainty in the estimated variation. The capacity study tracks method differences across $K = 1, \dots, 20$. The budget study tests whether more search changes the one-to-two abstraction loss. Both use maximum- t intervals from 10,000 seed-level bootstrap resamples for related estimates. Experiment I registered its contrasts and directions before data collection, but selected the Student- t interval afterward. The capacity analysis preceded fresh data. Its results motivated the registered budget hypothesis and budget increase. Later analyses are descriptive or post hoc.

4 Results

4.1 Experiment I: Which Abstractions?

With ten abstractions, validation utility solved 66.08 test problems (Figure 2). Validation-solution compression acquired solutions for a mean of 16.59 of the 25 validation problems and omitted 33.6%. It solved 62.00 test problems. The difference was 4.08 (95% CI [2.46, 5.71]). Both used one menu and validation set. Utility scored all outcomes, including failures. Compression scored acquired programs and omitted the rest. Thus, the difference includes the scoring objective and solution availability. The primitive-only library solved 58.91 problems. Standard compression, which used all 100 starter problems, solved 64.53. Validation-solution compression solved 0.26 fewer problems than compression on the 25-problem starter subset (95% CI [−1.46, 0.95]). Utility’s 1.55-problem advantage over standard compression remained uncertain because its 95% CI [−0.27, 3.37] included zero. Utility outperformed validation-solution compression. In the separate stale-validation stress test, utility solved a mean of 65.60 test problems. Validation-solution compression solved 61.00, and standard compression solved 65.40.

Library selection has an initial search cost. A more expensive selection method is useful only if later search savings recover that initial cost. Utility has a higher initial search cost than compression because it performs one bounded search for each trial library and scores all 25 validation targets from its output record. Selected abstractions can later save search because the solver can reuse stored grids instead of rebuilding them from primitive operations. We therefore estimate how many future problems are needed for these savings to recover Utility’s higher initial search cost. We call this the break-even point. The median break-even point was several thousand future problems. Some seed-condition cases never recovered Utility’s higher initial search cost. Figure 3 gives the exact selection costs, median break-even points, and case counts.




Figure 2: Experiment I held-out results with ten abstractions. The left graph shows means. The right graph shows paired differences with 95% Student t intervals for 30 seeds. Both procedures use the same validation problems. Utility includes failures; compression omits unsolved targets. The second contrast compares validation-solution compression with compression on 25 starter problems.




Figure 3: Selection cost and break-even estimates. Gray labels mark cases with no break-even point. Counts omit standard-compression acquisition, wall-clock time, and segmentation.

We expected compression to benefit from similar motifs in starter and test problems. As hypothesized, utility minus standard compression decreased as problem-set similarity increased. The mean difference was 2.7 problems for different sets, 1.5 for moderately similar sets, and 0.5 for similar sets (left graph in Figure 4). Assigned shift changed utility minus validation-solution compression by 1.70 problems (95% CI $[-0.56, 3.96]$). The seed-controlled estimate for problem-set similarity was -0.36 (95% CI $[-2.43, 1.71]$); the slope versus standard compression was -1.44 (95% CI $[-3.18, 0.30]$). The different-set interval excluded zero, showing utility's advantage in that group. Formal comparisons across problem-set similarity levels included zero and did not establish the hypothesized effect. We also tested whether abstractions chosen via search utility also indirectly compressed solutions. We applied each library post hoc to the same starter solutions (right graph in Figure 4). Direct compression on 25 starter problems removed 25.78% of operations; utility on those problems removed 12.32%; validation utility on 25 validation targets removed 9.79%. Utility shortened solutions indirectly. Direct compression removed more.




Figure 4: Problem-set similarity and indirect compression. Left: utility minus standard compression at three problem-set similarity levels. Right: post hoc operations removed from the same starter solutions. Intervals are 95% seed-cluster and seed bootstrap, respectively; problem-set similarity bands were not registered.

4.2 Experiment II: Library Size

Held-out performance did not change steadily as library size increased (Figure 5). From ten to twenty abstractions, standard compression lost 12.35 problems (simultaneous 95% CI $[-13.89, -10.81]$), and utility lost 12.11 (simultaneous 95% CI $[-13.20, -11.02]$). At twenty abstractions, both libraries solved fewer problems than the primitive-only library. Utility exceeded validation-solution compression by 5.13 problems (simultaneous 95% CI $[3.45, 6.80]$) and standard compression by 1.73 (simultaneous 95% CI $[-0.17, 3.63]$). The utility advantage over validation-solution compression increased by 1.85 from ten to twenty (simultaneous 95% CI $[-0.03, 3.73]$). Thus, library size had a larger effect on absolute performance than on method differences. At $B = 30,000$, the descriptive means were highest at $K = 11$: 66.32 for standard compression and 67.85 for utility. From $K = 1$ to $K = 2$, standard compression lost 8.08 problems, and utility lost 6.14. For standard compression, operations removed from starter solutions increased from 24.14% at $K = 11$ to 27.21% at $K = 20$. Held-out performance decreased over the same range. Thus, more compression did not ensure better bounded search. This result motivated Experiment III.

4.3 Experiment III: Search Budget

An increased budget reversed the one-to-two decrease for the fixed libraries (Figure 5). As the budget increased from 30,000 to 90,000 program attempts, the second-abstraction effect increased by 11.42 solved problems for standard compression. Its simultaneous 95% CI was $[10.17, 12.66]$. The effect increased by 9.61 problems for validation utility. Its interval was $[8.52, 10.69]$. At $B = 90,000$, the effects were 3.33 for standard compression and 3.47 for validation utility. All seed-level interactions were positive. With two abstractions, the share of enumerations that reached size-four programs increased from 8.1% at 30,000 attempts to 99.2% at 45,000 and 100% at 60,000. By $B = 90,000$, the solver recovered all test instances that the two-abstraction library lost relative to the one-abstraction library at $B = 30,000$. Only the budget changed. Thus, the 30,000-attempt budget caused the one-to-two decrease for these fixed libraries. This result does not set a sufficient budget for other libraries, solvers, or benchmarks.




Figure 5: Experiments II and III. The top row shows held-out performance and starter-solution compression across library sizes. The held-out bands are simultaneous 95% intervals. The compression graph is post hoc. At $K = 0$, the library has only primitives. The bottom row shows the fixed-library budget test, size-four access, and descriptive recovery of losses.

5 Discussion

The central conclusion is that an abstraction has no fixed value. Its value depends on selection evidence and objective, library contents, future problems, solver, search budget, selection cost, and expected reuse. Experiment I compares the complete selection procedures. Validation utility minimized capped search cost for all 25 validation targets, including failures. Validation-solution compression minimized operations in acquired solutions for the same targets but omitted unsolved targets. With ten abstractions, utility solved 4.08 more held-out test problems than validation-solution compression. This supports the hypothesis that complete validation utility selects more useful abstractions than validation-solution compression. Utility measures bounded search cost, including added programs and full-budget failures. Compression measures shortcuts only in acquired solutions. Utility can thus favor lower net search cost. However, the 4.08-problem difference also reflects solution availability, so it does not isolate the scoring objective. Utility’s 1.55-problem advantage over standard compression remained uncertain. Standard compression used 100 starter problems, four times as many. Broader evidence could explain the smaller difference, but this experiment did not test that explanation.

How well selection problems match future problems could also change abstraction value. We therefore hypothesized that compression would become more competitive as starter and test motif frequencies became more similar. The means followed this pattern: utility’s advantage over standard compression fell from 2.7 problems for different sets to 0.5 for similar sets. However, the formal intervals included zero, so the statistical evidence remains inconclusive. This direction follows be-

cause compression rewards stored grids that replace subtrees in starter solutions. Recurring motifs can let those grids shorten test solutions. Our problem-set similarity measure, however, ranks only motif counts and omits operator context, exact grid reuse, program position, and search order. Similar motif counts do not ensure similar search value.

Experiment II showed that abstraction value depends on library size. At 30,000 attempts, observed mean performance was highest at 11 abstractions and fell below the primitive-only library at 20. Standard compression removed more operations from starter solutions as its library grew from 11 to 20 abstractions, even as held-out performance decreased. Each abstraction is a size-zero item. Combining it with operators and other library items creates new low-operation programs. The solver examines them first, so they can use the budget before it reaches a larger program where the abstraction helps. A new item also changes when the solver reaches programs that use earlier items. Thus, operation removal measures a shortcut within known solutions. It does not measure the search required to reach those solutions.

Search budget can determine whether this added search prevents a solution or only delays it. Experiment III changed only the budget for fixed one- and two-abstraction libraries. The larger budget reversed the one-to-two loss for both procedures. This supports the budget hypothesis for those libraries. The share of enumerations that reached size-four programs rose from 8.1% at 30,000 attempts to 99.2% at 45,000. This diagnostic gives descriptive support for the search-cutoff explanation. However, the experiment does not define a sufficient budget for other library sizes, solvers, or tasks.

Selection cost creates a separate tradeoff because future savings must recover the initial cost. Validation utility required a bounded search for every trial library. Among cases with a break-even point, the median required thousands of future problems. Some experimental cases never reached that point. Existing representative solutions reduce compression’s initial selection cost. The break-even calculation counts only program attempts, so it does not estimate total deployment cost.

6 Limitations and Future Work

The controls make changes in bounded search easier to identify, but they limit where the results can apply. This study tests one synthetic, target-only grid benchmark. Every abstraction is a fixed output grid with size zero. The primitives, operators, program-size limit, and target generator stay fixed. Learned functions and parameterized abstractions can apply across inputs and produce different search costs. Future work should repeat these experiments on input–output synthesis tasks and in other domains. A direct comparison of stored grids and learned functions should hold the solver and budget fixed.

The search-cutoff explanation also depends on the solver. This solver performs deterministic bottom-up enumeration and orders programs by operation count. Library order and duplicate removal affect when the solver reaches a target. A guided solver could prune low-value combinations or reach useful larger programs sooner. Future studies should compare deterministic enumeration with guided search and test other costs for abstraction use. A joint sweep can test whether library size and test performance have the same relation under other budgets.

Candidate discovery also limits the conclusions. All methods select from the same 50 starter-derived grids. Validation utility cannot select an abstraction outside this menu. The menu favors grids that support at least two starter targets and first appear in programs with one to four operations. Experiment II evaluates prefixes of one greedy sequence rather than optimizing each library size. Future work should vary menu construction and compare greedy selection with joint selection of library items and their order. A selection objective could also charge for the new programs that each candidate adds.

The primary Experiment I comparison evaluates complete procedures, so it does not isolate the scoring objective. Utility scored all 25 validation targets and failures; compression scored only acquired solutions. The 4.08-problem gap therefore includes the scoring objective and solution availabil-

ity. Standard compression used solutions acquired from 100 starter problems, so its evidence also differed. A same-evidence comparison should use fixed target-solution pairs. Utility would score the targets, and compression would score their solutions. Solution acquisition must be fixed before data collection. The similarity measure has a separate limit: it requires hidden motif counts from completed test data, so it cannot guide deployment. Future work should define the measure before data collection and vary compositional structure directly. More independent seeds would improve precision.

Several findings need independent confirmation. The observed performance peak near 11 abstractions and the one-to-two decrease were descriptive. Experiment III reused the Experiment II seeds and fixed one- and two-abstraction libraries. Its budget effect applies only to those libraries and is not an independent replication. An independent test should use new seeds and libraries. A broader sweep can identify the budgets needed to reach programs that use added abstractions.

Finally, the break-even analysis counts only program attempts. It omits wall-clock time, memory use, and compression segmentation. The standard-compression comparison treats prior starter-solution acquisition as sunk. The calculation assumes constant mean savings across future problems. Future work should measure end-to-end acquisition, selection, and evaluation costs for expected workloads under distribution change.

7 Conclusion

An abstraction has no fixed value. Its value depends on the data and score used for selection, the future problems, the library, and the solver. With ten abstractions, validation utility solved 4.08 more test problems than validation-solution compression. This difference compares complete procedures and includes the scoring objective and solution availability. Utility’s advantage over standard compression remained uncertain. Mean performance differences across problem-set similarity levels followed the hypothesis. Formal intervals for these differences included zero.

Library size also changed whether selected abstractions helped. At 30,000 attempts, both procedures improved on primitive-only search with ten abstractions. Mean performance peaked near 11 and fell below primitives alone at 20. Standard compression removed more operations from starter solutions over this range. Held-out performance decreased. Solution compression alone was an incomplete measure of bounded-search value. Each added item made more low-operation programs available. The solver examined these before later programs where the item could help.

For the fixed libraries, search budget determined whether the solver reached those later programs. Increasing only the budget reversed the one-to-two loss. This result supports the search-cutoff explanation for those libraries. Utility’s added selection work often required thousands of future problems to recover and did not recover in some cells. A practical system should evaluate the complete library with its intended solver, budget, workload, and expected reuse. A useful abstraction must save more search than its selection and added programs cost on that workload.

Code and Data Availability

The public repository contains the code, generated data, and reproduction instructions.

github.com/cucupac/program-synthesis.

References

- Matthew Bowers, Theo X. Olausson, Lionel Wong, Gabriel Grand, Joshua B. Tenenbaum, Kevin Ellis, and Armando Solar-Lezama. Top-down synthesis for library learning. *Proceedings of the ACM on Programming Languages*, 7(POPL):1182–1213, 2023. DOI: 10.1145/3571234.
- Kevin Ellis, Catherine Wong, Maxwell Nye, Mathias Sablé-Meyer, Lucas Morales, Luke Hewitt, Luc Cary, Armando Solar-Lezama, and Joshua B. Tenenbaum. DreamCoder: Bootstrapping inductive program synthesis with wake-sleep library learning. In *Proceedings of the 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation*, pp. 835–850. Association for Computing Machinery, 2021. DOI: 10.1145/3453483.3454080.
- Leonardo Hernandez Cano, Ivan Zareski, Luisa El Amouri, Pinzhe Zhao, Max Mascini, Emanuele Sansone, Yewen Pu, Bonan Zhao, and Marta Kryven. Prospective compression in human abstraction learning. *arXiv*, 2026. DOI: 10.48550/arXiv.2605.09985.
- Glenn A. Iba. A heuristic approach to the discovery of macro-operators. *Machine Learning*, 3(4): 285–317, 1989. DOI: 10.1007/BF00116836.
- Zhentao Ye, Ruyi Ji, Yingfei Xiong, and Xin Zhang. Accelerating syntax-guided program synthesis by optimizing domain-specific languages. *Proceedings of the ACM on Programming Languages*, 10(POPL):1066–1093, 2026. DOI: 10.1145/3776679.
- Janis Zenkner, Lukas Dierkes, Tobias Sesterhenn, and Christian Bartelt. AbstractBeam: Enhancing bottom-up program synthesis using library learning. *arXiv*, 2024. DOI: 10.48550/arXiv.2405.17514.